

# Lecture 3

# Convolutional Neural Networks

*Minh-Tan Pham, Matthieu Le Lain*

**BUT2 INFO, 2023-2024**

**IUT de Vannes, Université Bretagne Sud**

*minh-tan.pham@univ-ubs.fr*

# Planning


- Why CNNs ?
- Image Convolution
- Pooling Layer
- Typical CNNs
- Lab Assignments

# Why CNNs ?

# Why CNNs ?

## Neural nets - Limitations


- Neural networks = fully-connected networks
- Ignore the input data structure (spatial relationship and topology)
- Require too many parameters  $\rightarrow$  more time, disk space, more training data



# Why CNNs ?

## Neural nets - Limitations


- Neural networks = fully-connected networks
- Ignore the input data structure (spatial relationship and topology)
- Require too many parameters → more time, disk space, more training data



# Why CNNs ?

## Solution: local connectivity


- Sparse connectivity → weights only connected to local patch



# Why CNNs ?

## Solution: local connectivity


- Sparse connectivity → weights only connected to local patch
- Preserve the spatial topology (keep spatial information in 2D)



# Why CNNs ?

## Solution: local connectivity

- Sparse connectivity → weights only connected to local patch
- Preserve the spatial topology (keep spatial information in 2D)
- Share the same parameters across different locations → reduce the number of parameters






# Why CNNs ?

## Convolutional layer


- Connect each hidden unit (neuron) to a small input patch (locally) → filter or kernel



# Why CNNs ?

## Convolutional layer


- Connect each hidden unit (neuron) to a small input patch (locally) → filter or kernel
- Share the same weights across spatial space



# Why CNNs ?


## Convolutional layer

- Connect each hidden unit (neuron) to a small input patch (locally) → filter or kernel
- Share the same weights across spatial space
- Learn multiple filters



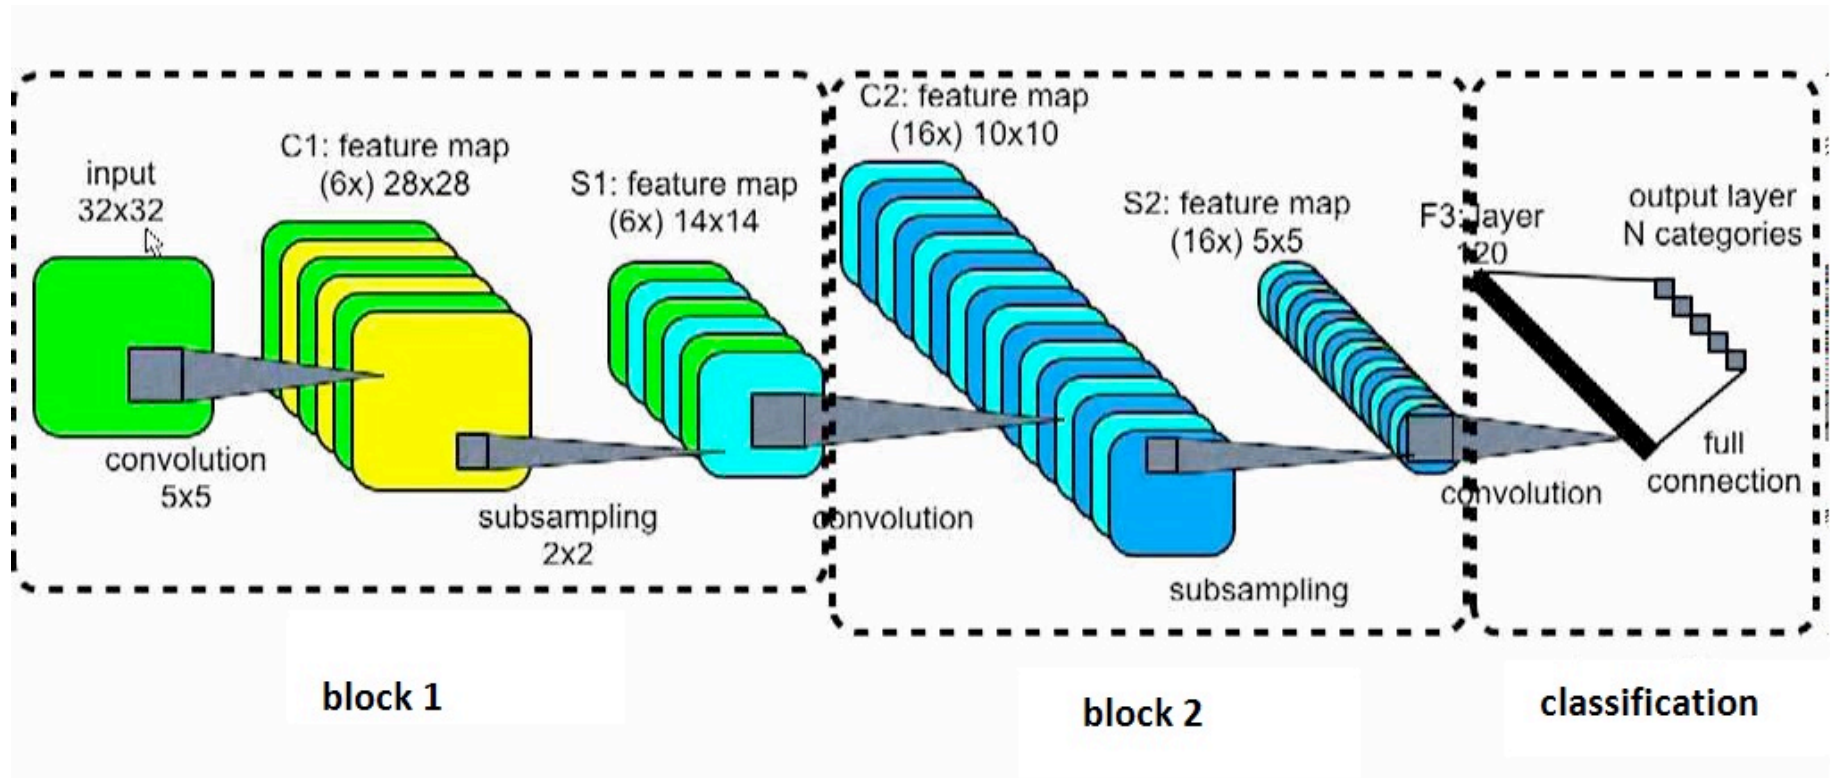
# Convolutional neural networks

- Replace the fully-connected layer by the convolutional layer
- Combine with other layers: non-linear activation, pooling (sub-sampling), normalization, drop out, etc. ) → an elementary block
- For example: a common used block = convolution + ReLu + Pooling



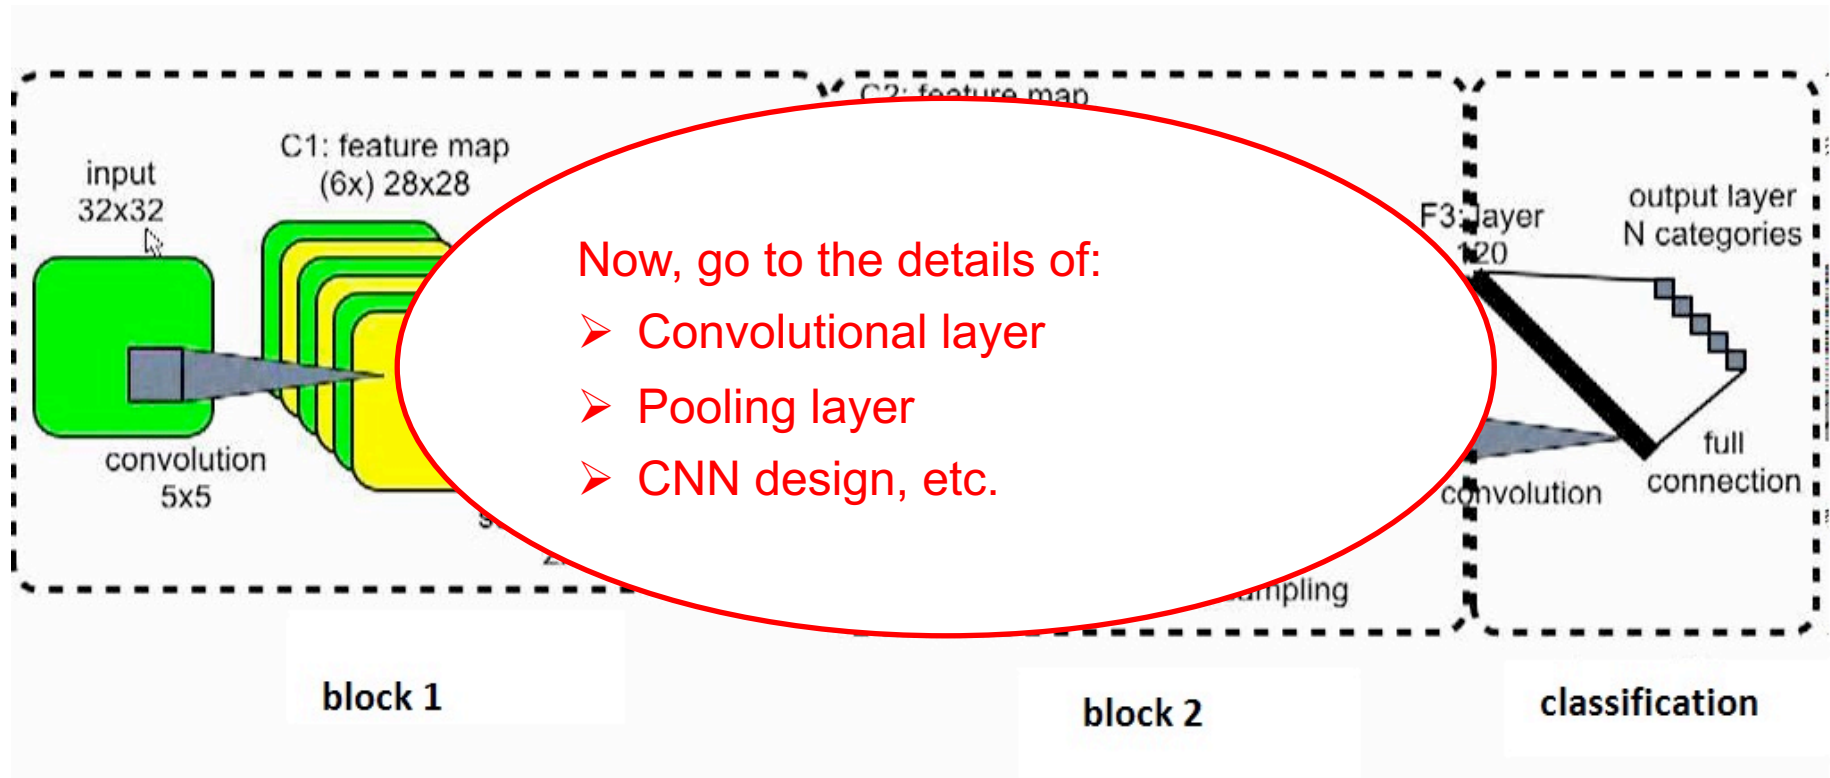
# Convolutional neural networks

- Replace the fully-connected layer by the convolutional layer
- Combine with other layers: non-linear activation, pooling (sub-sampling), normalization, drop out, etc. ) → an elementary block
- For example: a common used block = convolution + ReLu + Pooling
- Stack several block → **CNNs**



# Convolutional neural networks

- Replace the fully-connected layer by the convolutional layer
- Combine with other layers: non-linear activation, pooling (sub-sampling), normalization, drop out, etc. ) → an elementary block
- For example: a common used block = convolution + ReLu + Pooling
- Stack several block → **CNNs**





# Convolutional neural networks

- Before we go, let's imagine !!!!

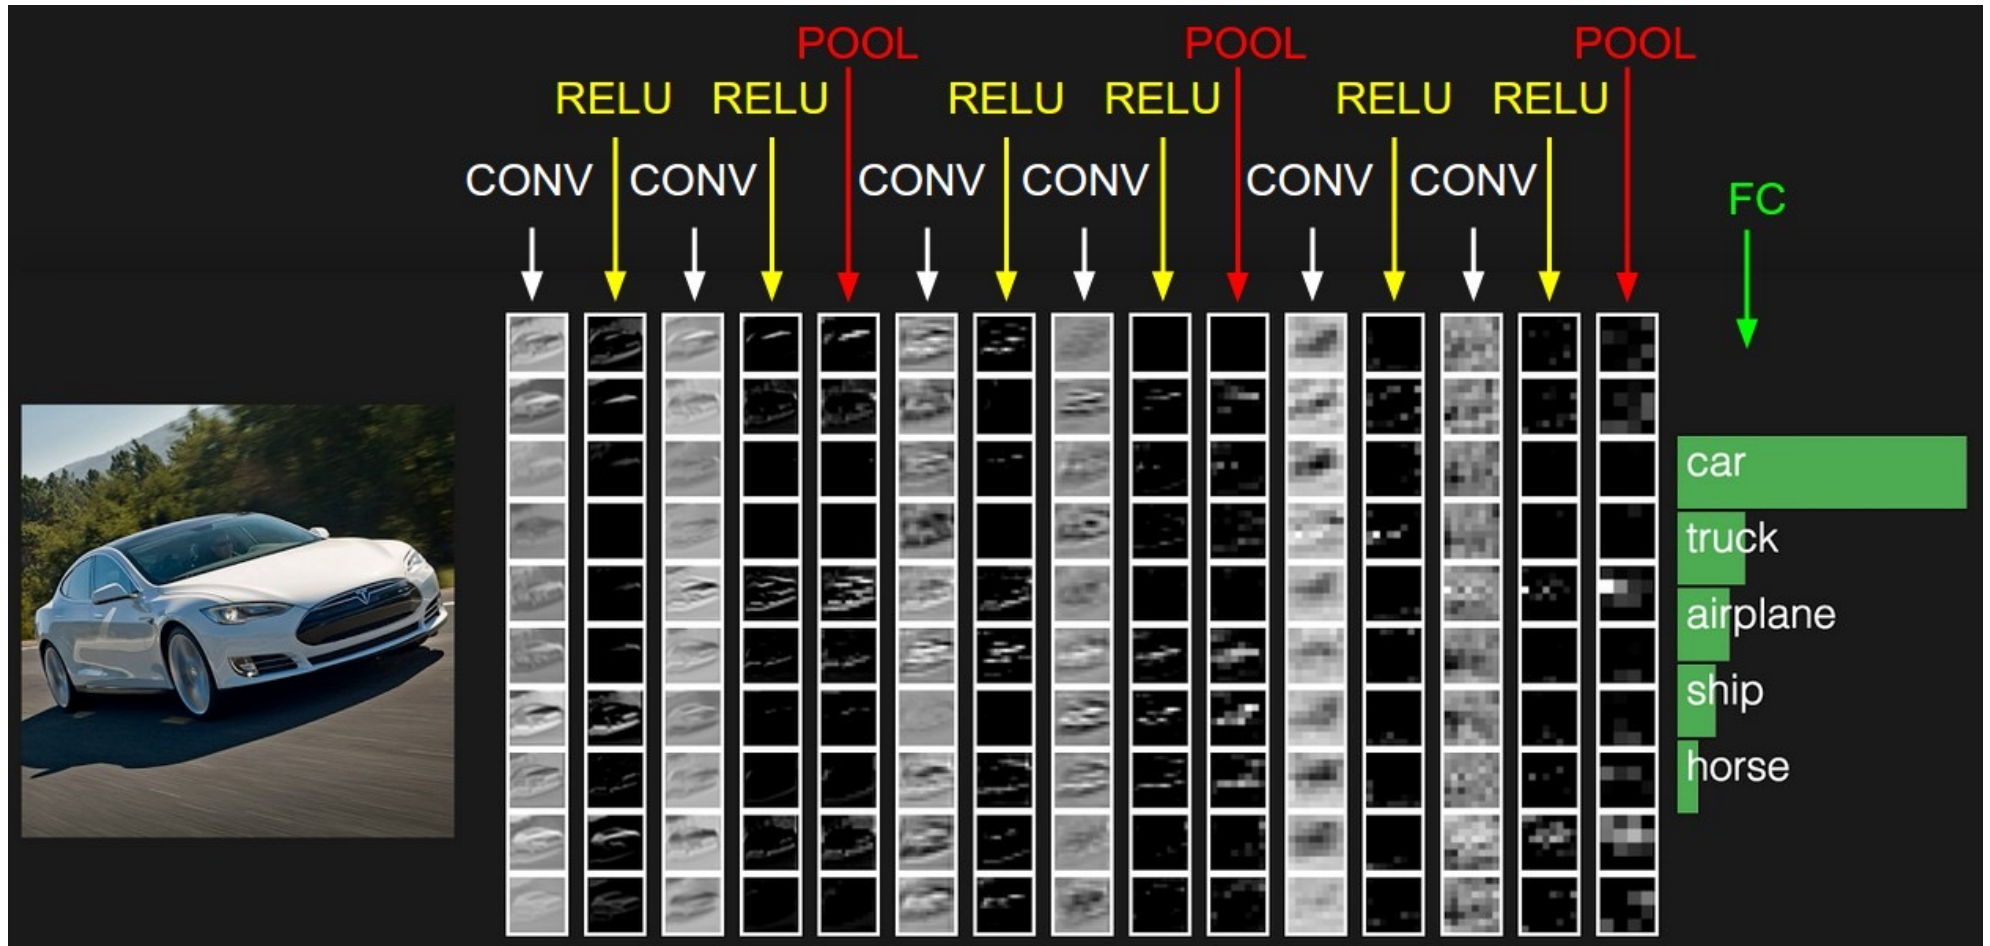


Image credits: [cs231n](#)

# Image Convolution



# Image Convolution - Grayscale image


- 2D convolution on image **f** with a filter **h** of size  $d \times d$

$$f'(i,j) = (f \star h)(i,j) = \sum_{n=-\frac{d-1}{2}}^{\frac{d-1}{2}} \sum_{m=-\frac{d-1}{2}}^{\frac{d-1}{2}} f(i-n, m-j) h(n, m)$$


- In practice:

→ Center the filter **h** at each pixel location

→ Compute the weighted sum



# Image Convolution - Grayscale image



# Image Convolution - Example

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Image

1	0	1
0	1	0
1	0	1

Filter/kernel

# Image Convolution - Example

1 <sub>x1</sub>	1 <sub>x0</sub>	1 <sub>x1</sub>	0	0
0 <sub>x0</sub>	1 <sub>x1</sub>	1 <sub>x0</sub>	1	0
0 <sub>x1</sub>	0 <sub>x0</sub>	1 <sub>x1</sub>	1	1
0	0	1	1	0
0	1	1	0	0

4		

# Image Convolution - Example

1	1 <sub>x1</sub>	1 <sub>x0</sub>	0 <sub>x1</sub>	0
0	1 <sub>x0</sub>	1 <sub>x1</sub>	1 <sub>x0</sub>	0
0	0 <sub>x1</sub>	1 <sub>x0</sub>	1 <sub>x1</sub>	1
0	0	1	1	0
0	1	1	0	0

4	3	

# Image Convolution - Example

1	1	1 <sub>x1</sub>	0 <sub>x0</sub>	0 <sub>x1</sub>
0	1	1 <sub>x0</sub>	1 <sub>x1</sub>	0 <sub>x0</sub>
0	0	1 <sub>x1</sub>	1 <sub>x0</sub>	1 <sub>x1</sub>
0	0	1	1	0
0	1	1	0	0

4	3	4

# Image Convolution - Example

1	1	1	0	0
0 <sub>x1</sub>	1 <sub>x0</sub>	1 <sub>x1</sub>	1	0
0 <sub>x0</sub>	0 <sub>x1</sub>	1 <sub>x0</sub>	1	1
0 <sub>x1</sub>	0 <sub>x0</sub>	1 <sub>x1</sub>	1	0
0	1	1	0	0

4	3	4
2		


# Image Convolution - Example

1	1	1	0	0
0	1	1	1	0
0	0	1 <sub>x1</sub>	1 <sub>x0</sub>	1 <sub>x1</sub>
0	0	1 <sub>x0</sub>	1 <sub>x1</sub>	0 <sub>x0</sub>
0	1	1 <sub>x1</sub>	0 <sub>x0</sub>	0 <sub>x1</sub>

4	3	4
2	4	3
2	3	4



# Image Convolution - Example



In general:

$$W_{f'} = W_f - W_h + 1$$

$$H_{f'} = H_f - H_h + 1$$

# Image Convolution - padding

1	1	2	5	6	3	6	7	3
2	3	4	6	7	5	1	8	4
8	7	6	5	7	6	3	3	4
2	3	5	6	7	8	2	7	3
4	5	3	2	1	6	8	7	2
1	4	5	3	2	6	7	8	1
2	3	4	5	6	8	9	2	1

Input Image

$\ast \frac{1}{25}$

1	1	1	1	1
1	1	1	1	1
1	1	1	1	1
1	1	1	1	1
1	1	1	1	1

filter



0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	1	1	2	5	6	3	6	7	3	0	0
0	0	2	3	4	6	7	5	1	8	4	0	0
0	0	8	7	6	5	7	6	3	3	4	0	0
0	0	2	3	5	6	7	8	2	7	3	0	0
0	0	4	5	3	2	1	6	8	7	2	0	0
0	0	1	4	5	3	2	6	7	8	1	0	0
0	0	2	3	4	5	6	8	9	2	1	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0

Zero Padding(option1)


2	2	2	3	4	6	7	5	1	8	4	4	4
1	1	1	1	2	5	6	3	6	7	3	3	3
1	1	1	1	2	5	6	3	6	7	3	3	7
3	2	2	3	4	6	7	5	1	8	4	4	8
7	8	8	7	6	5	7	6	3	3	4	4	3
3	2	2	3	5	6	7	8	2	7	3	3	7
5	4	4	5	3	2	1	6	8	7	2	2	7
4	1	1	4	5	3	2	6	7	8	1	1	8
3	2	2	3	4	5	6	8	9	2	1	1	2
3	2	2	3	4	5	6	8	9	2	1	1	2
3	1	1	4	5	3	2	6	7	8	1	1	2

Repetition of border pixels(option 2)

→ Output has the same size as input !


# Image Convolution - stride

Example : stride of 2




# Image Convolution – output size

$$W_{output} = \frac{(W_{input} + 2 \times padding) - W_{filter}}{stride} + 1$$




# Image Convolution - output size

$$W_{output} = \frac{(W_{input} + 2 \times padding) - W_{filter}}{stride} + 1$$




# Image Convolution - output size

$$W_{output} = \frac{(W_{input} + 2 \times padding) - W_{filter}}{stride} + 1$$




# Convolution: multi-channel image

- Each filter has the size of  $d \times d \times D$  (spatial and depth)
- Common color image:  $D=3$  (3 bands RGB)
- Example of multispectral image:  $D=4$  (RGB+IR)



# Convolution: multi-channel image

- Each filter has the size of  $d \times d \times D$  (spatial and depth)
- Common color image:  $D=3$  (3 bands RGB)
- Example of multispectral image:  $D=4$  (RGB+IR)
- Use multiple filters  $\rightarrow$  multiple output filtered images






# Pooling Layer

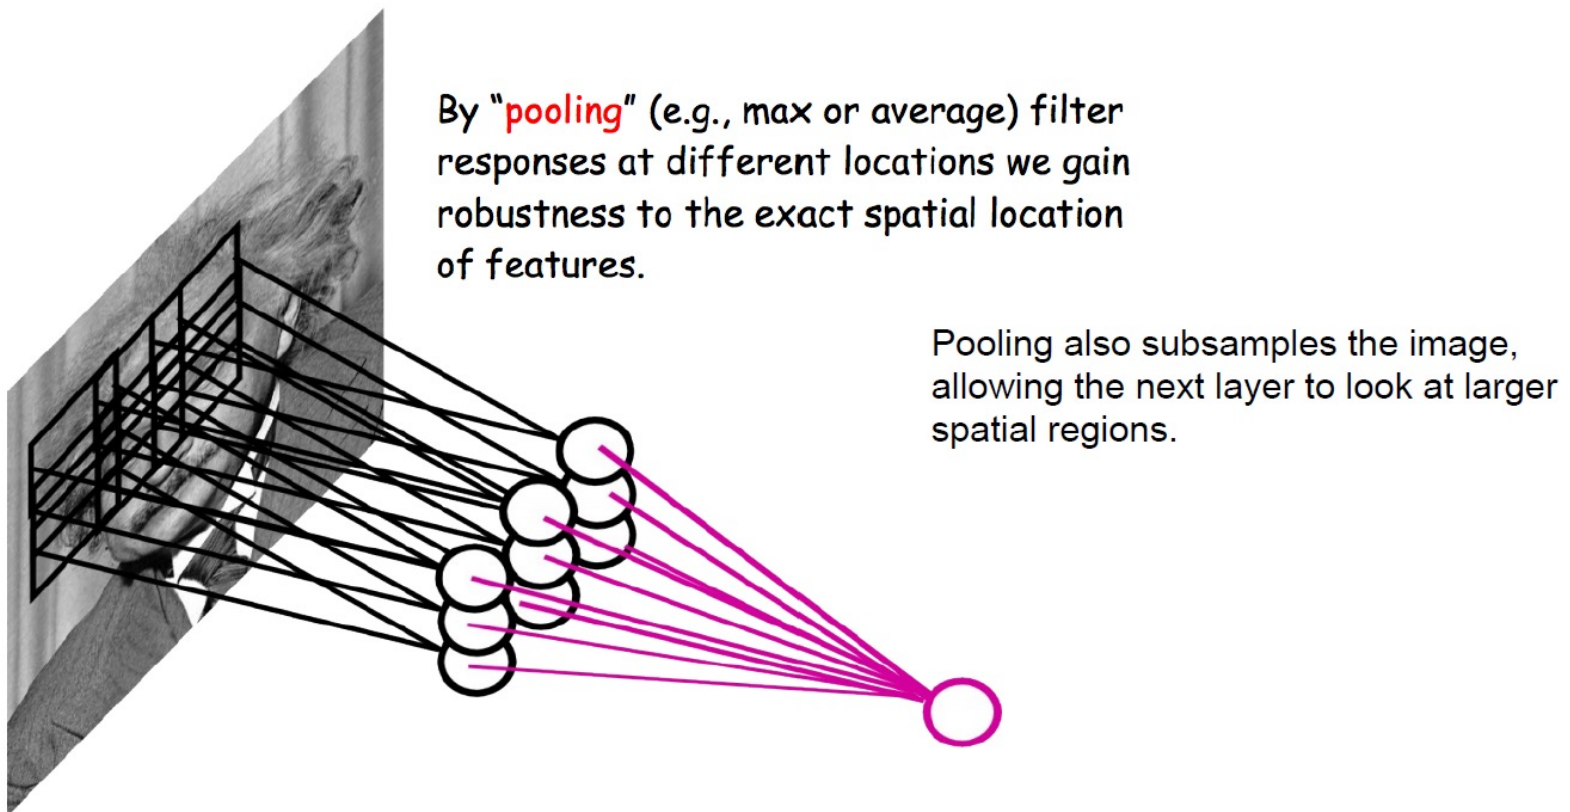
# Pooling layer

- spatial down-sampling (sub-sampling) ! reduce spatial size
- reduce number of parameters and computation time
- increase shift/translation invariance
- average pooling or max-pooling (mostly used)



# Pooling layer

- spatial down-sampling (sub-sampling) ! reduce spatial size
- reduce number of parameters and computation time
- increase shift/translation invariance
- average pooling or max-pooling (mostly used)



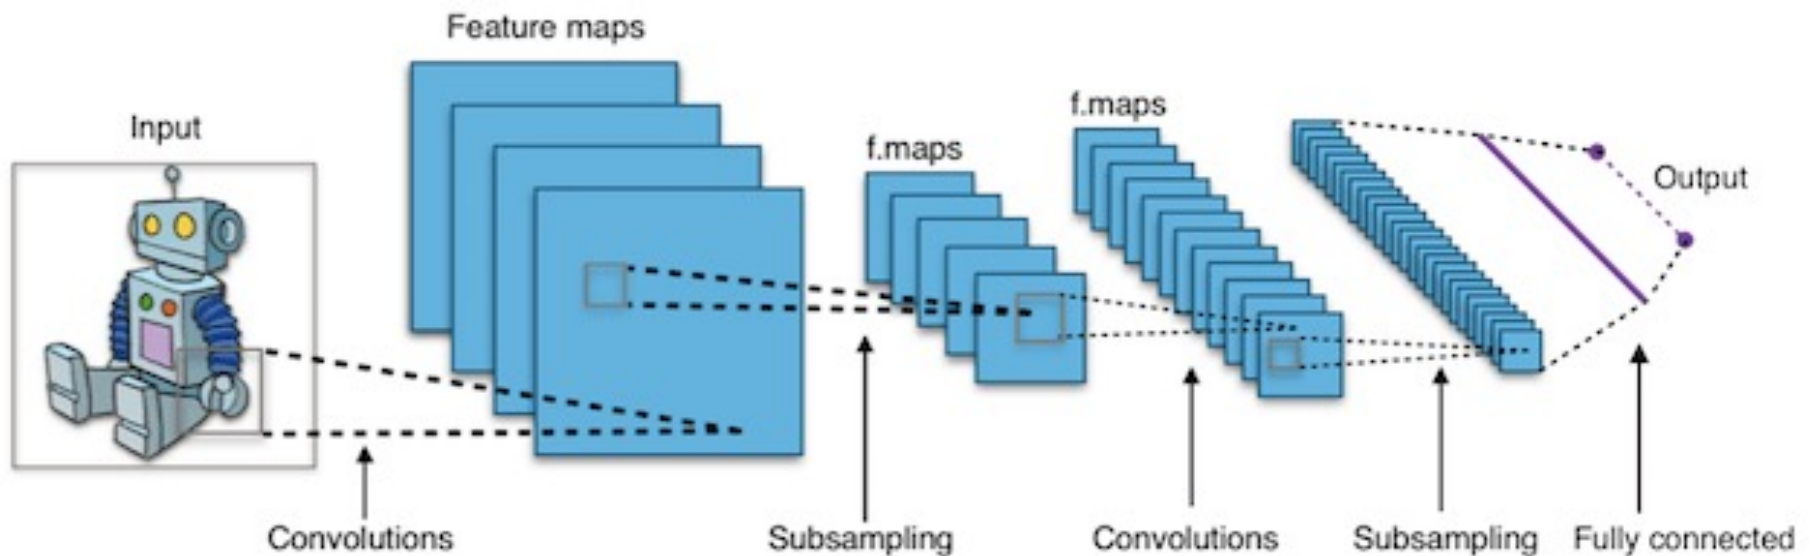
# Typical CNNs

# Typical CNN architecture

- A typical CNN architecture includes:

**Input image → [Conv - ReLu] → (Pool) → [Conv - ReLu] → (Pool)  
→ ... → FC (fully-connected) → Softmax (classification) → Output**

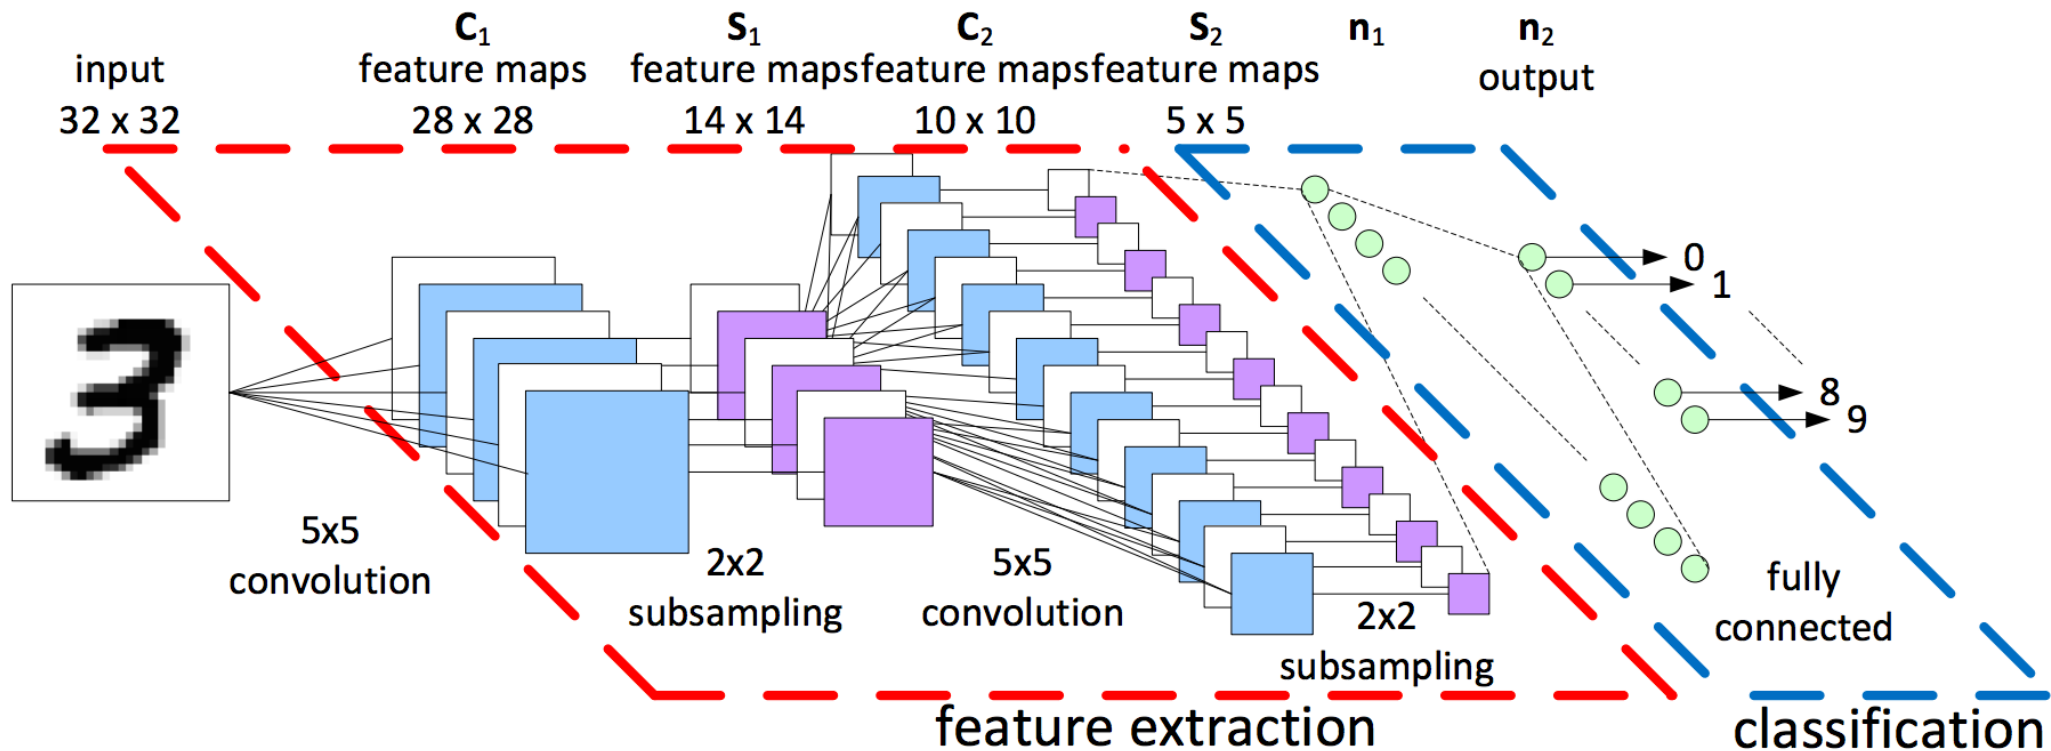
- Optional (for practical use): stride, normalization, dropout, etc.



# Typical CNN architecture

- A typical CNN architecture includes:

Input image  $\rightarrow$  [Conv - ReLu]  $\rightarrow$  (Pool)  $\rightarrow$  [Conv - ReLu]  $\rightarrow$  (Pool)  
 $\rightarrow$  ...  $\rightarrow$  FC (fully-connected)  $\rightarrow$  Softmax (classification)  $\rightarrow$  Output



# Typical CNN architecture

- Let look back to the example


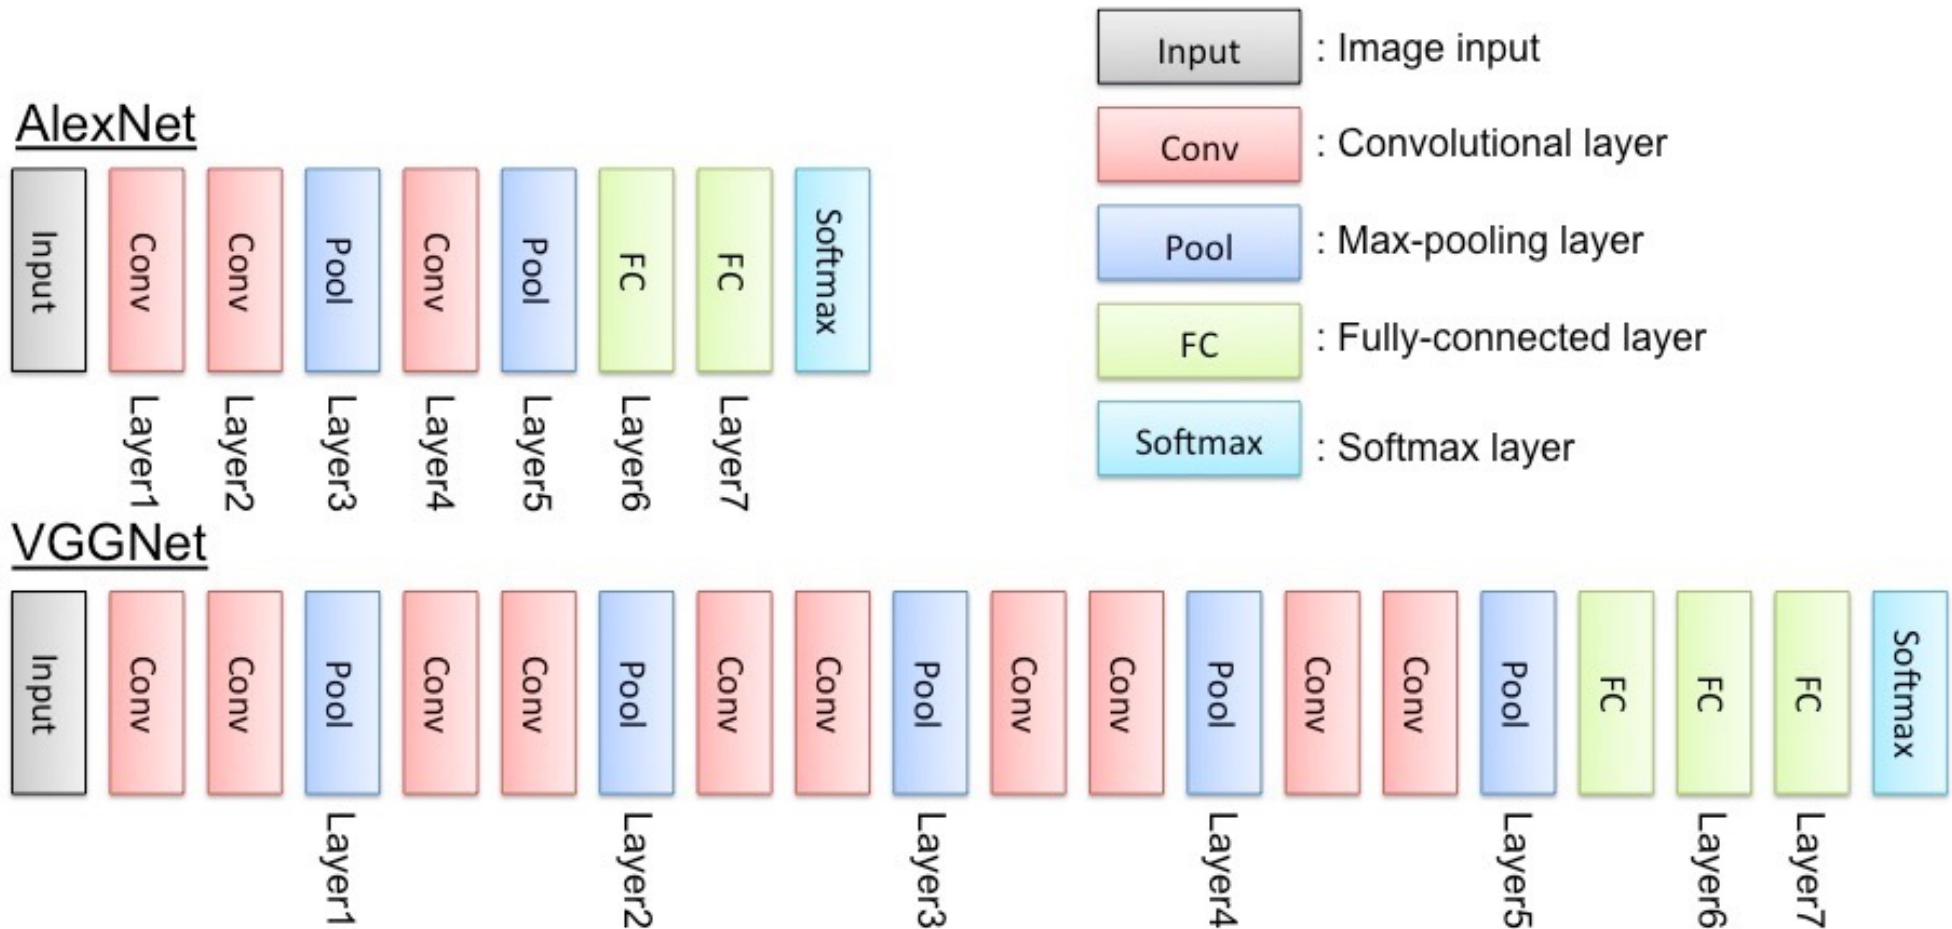


Image credits: [cs231n](https://www.cs231n.org/)



# Typical CNN architecture





# Typical CNN architecture

AlexNet



# Lab assignment

**Practical lab:** Training convolutional neural networks for digit recognition (MNIST) and for color object detection (CIFAR10) using Pytorch

- Download, complete and submit the notebook from Moodle (**R3A13\_lab2\_assign.ipynb**)
- Note that you need to submit **your compiled notebook with all outputs**

# References and Sources

- **Convolutional neural networks for visual recognition** - Stanford

<https://cs231n.github.io/>

- **Neural networks and deep learning** (free online book)

<http://neuralnetworksanddeeplearning.com/index.html>

- **Machine learning course** – Oxford

<https://www.cs.ox.ac.uk/people/nando.defreitas/machinelearning/>

- **Neural network course** - Hugo Larochelle

[https://info.usherbrooke.ca/hlarochelle/neural\\_networks/content.html](https://info.usherbrooke.ca/hlarochelle/neural_networks/content.html)

- **Deep learning course** – François Fleuret

<https://fleuret.org/dlc/>